

TEDx

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import seaborn as sns; sns.set()
from urllib.request import urlretrieve
from IPython.display import clear_output
import cv2
from PIL import Image
from sklearn.metrics.cluster import entropy as sklearn_entropy
from skimage import io, color, img_as_ubyte
from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
import string
from textblob import TextBlob
import traceback
import json
import requests
import time
from bs4 import BeautifulSoup
```

Let's check how the DF looks

```
In [2]: df = pd.read_excel("data/data_tedx.xlsx")
df.head()
```

Out[2]:

	Licensees.name	Events.name	Videos.id	Videos.videoid	Videos.title	
0	TEDxRíodelaPlata	CanalTEDxRíodelaPlata	72321a94-f58e-47a2-9071-40a5ae54091e	qJzRzdsGhC0	Flash Mob en TEDxRíodelaPlata 2013 - El Brindi...	ht /
1	TEDxRíodelaPlata	TEDxRíodelaPlata2013	ea1705b1-2eb7-4782-a9d3-4f88fcc0de5a	EFdEmbuikOw	Nunca pidas permiso Ronald Shakespear TEDx...	ht /v
2	TEDxRíodelaPlata	TEDxJoven@RíodelaPlata 2011	54489a1b-12e8-4e5a-b9aa-bba37452309d	nZV1uefYywl	Sobre fideos y jabones: Luciano Mellera at TED...	ht
3	TEDxRíodelaPlata	TEDxRíodelaPlata2014	e5a40182-5cec-4e28-afe6-784750a83cd3	jej8qlzIAGw	Las matematicas son para siempre Eduardo Sae...	ht
4	TEDxRíodelaPlata	TEDxRíodelaPlata2015	ea927c7b-c454-4fa4-ae54-86900daa0fff	c1iMoel2LxA	Secretos de la relación corazón-cerebro Luci...	ht

```
In [3]: df.columns
```

```
Out[3]: Index(['Licensees.name', 'Events.name', 'Videos.id', 'Videos.videoid',
               'Videos.title', 'Videos.url', 'Videos.url_img', 'Videos.description',
               'Videos.published_at', 'Videos.channelid', 'Videos.channel_title',
               'Videos.tags', 'Videos.view_count', 'Videos.like_count',
               'Videos.dislike_count', 'Videos.comment_count', 'Videos.users_id',
               'Videos.licensees_id', 'Videos.created', 'Videos.modified'],
              dtype='object')
```

Unique values

```
In [4]: df.nunique()
```

```
Out[4]: Licensees.name          1
        Events.name           15
        Videos.id            265
        Videos.video_id      265
        Videos.title         265
        Videos.url           265
        Videos.url_img       265
        Videos.description    257
        Videos.published_at   259
        Videos.channel_id     3
        Videos.channel_title  3
        Videos.tags           240
        Videos.view_count     265
        Videos.like_count     217
        Videos.dislike_count  82
        Videos.comment_count  97
        Videos.users_id       5
        Videos.licensees_id   0
        Videos.created        183
        Videos.modified       7
        dtype: int64
```

```
In [5]: df = df.drop(['Licensees.name', 'Videos.licensees_id'], axis=1)
```

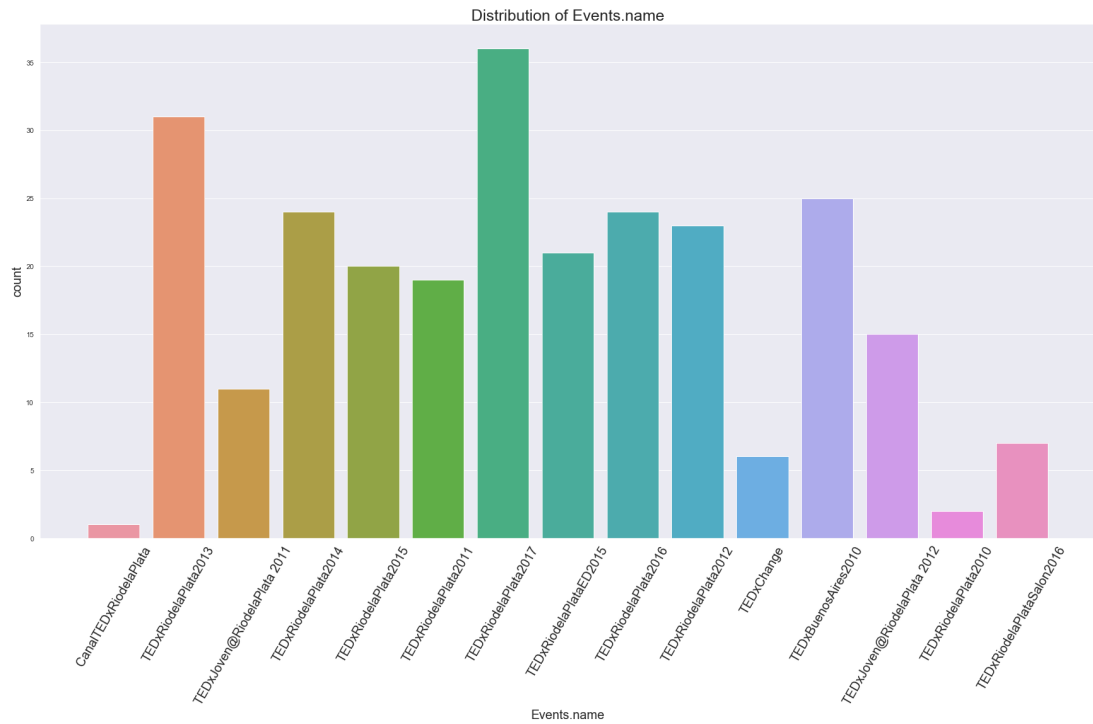
```
In [6]: df['Events.name'].unique()
```

```
Out[6]: array(['CanalTEDxRiodelaPlata', 'TEDxRiodelaPlata2013',
               'TEDxJoven@RiodelaPlata 2011', 'TEDxRiodelaPlata2014',
               'TEDxRiodelaPlata2015', 'TEDxRiodelaPlata2011',
               'TEDxRiodelaPlata2017', 'TEDxRiodelaPlataED2015',
               'TEDxRiodelaPlata2016', 'TEDxRiodelaPlata2012', 'TEDxChange',
               'TEDxBuenosAires2010', 'TEDxJoven@RiodelaPlata 2012',
               'TEDxRiodelaPlata2010', 'TEDxRiodelaPlataSalon2016'], dtype=object)
```

```
In [7]: def graph_categorical_dist(feature, rotation=False):
        fig, ax = plt.subplots(figsize=(30,15))
        sns.countplot(data=df, x=feature)
        plt.xlabel(feature, fontsize=20)
        plt.ylabel('count', fontsize=20)
        plt.title('Distribution of %s' % feature, fontsize=25)
        for tick in ax.xaxis.get_major_ticks():
            tick.label.set_fontsize(20)
            if rotation:
                tick.label.set_rotation(rotation)
        sns.despine()
        plt.plot();
```

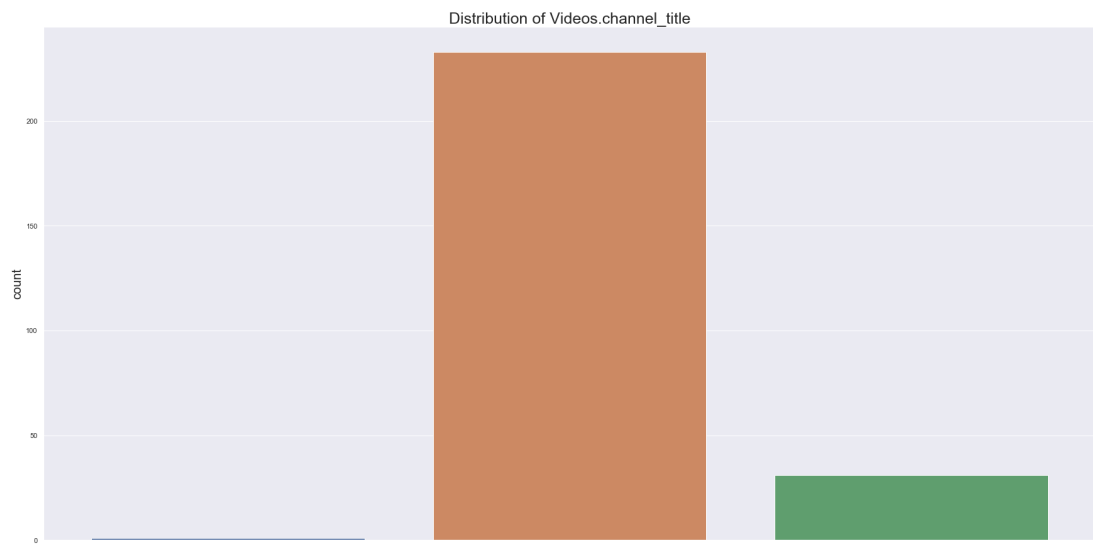
Let's check the distribution of events

```
In [8]: graph_categorical_dist('Events.name', rotation=60)
```



Let's check distribution of channels

```
In [9]: graph_categorical_dist('Videos.channel_title')
```



```
In [10]: df = df.drop(['Videos.videoid', 'Videos.url', 'Videos.description', 'Videos.users_id
```

```
In [11]: df.columns = ['event_name', 'video_id', 'title', 'url_thumb',
                        'tags', 'view_count', 'like_count',
                        'dislike_count', 'comment_count', 'created']
df.head()
```

Out[11]:

	event_name	video_id	title	url_thumb	
0	CanalTEDxRiodelaPlata	72321a94-f58e-47a2-9071-40a5ae54091e	Flash Mob en TEDxRiodelaPlata 2013 - El Brindi...	https://i.ytimg.com/vi/qJzRzdsGhC0/hqdefault.jpg	
1	TEDxRiodelaPlata2013	ea1705b1-2eb7-4782-a9d3-4f88fcc0de5a	Nunca pidas permiso Ronald Shakespear TEDx...	https://i.ytimg.com/vi/EFdEmbuikOw/hqdefault.jpg	Diseño, talks
2	TEDxJoven@RiodelaPlata 2011	54489a1b-12e8-4e5a-b9aa-bba37452309d	Sobre fideos y jabones: Luciano Meller at TED...	https://i.ytimg.com/vi/nZV1uefYywl/hqdefault.jpg	Luciano, Meller
3	TEDxRiodelaPlata2014	e5a40182-5cec-4e28-afe6-784750a83cd3	Las matematicas son para siempre Eduardo Sae...	https://i.ytimg.com/vi/jej8qlzLAGw/hqdefault.jpg	(hard), Cultur
4	TEDxRiodelaPlata2015	ea927c7b-c454-4fa4-ae54-86900daa0fff	Secretos de la relación corazón-cerebro Luci...	https://i.ytimg.com/vi/c1iMoel2LxA/hqdefault.jpg	TEDxTalks, Spanish

From "Engagement and Popularity Dynamics of YouTube Videos and Sensitivity to Meta-Data": Meta features from thumbnails: blurriness (e.g., CannyEdge, Laplace Frequency), brightness, contrast (e.g., tone), overexposure, and entropy of the thumbnail.
 blurriness: <https://www.pyimagesearch.com/2015/09/07/blur-detection-with-opencv/>
[\(https://www.pyimagesearch.com/2015/09/07/blur-detection-with-opencv/\)](https://www.pyimagesearch.com/2015/09/07/blur-detection-with-opencv/)

Saturacion, profundidad de campo, balance de blancos, r, g y b, ruido,

```
In [12]: df['blurriness'] = 0.0
df['brightness'] = 0.0
df['contrast'] = 0.0 # HOW?!
df['overexposure'] = 0.0 # HOW?
df['entropy'] = 0.0
df['avg_color'] = np.empty((len(df), 0)).tolist()
```

```
In [13]: def blurriness(path):
# compute the Laplacian of the image and then return the focus
# measure, which is simply the variance of the Laplacian
image = cv2.imread(path)
gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
return cv2.Laplacian(image, cv2.CV_64F).var()

def brightness(path):
image = Image.open(path)
greyscale_image = image.convert('L')
histogram = greyscale_image.histogram()
pixels = sum(histogram)
brightness = scale = len(histogram)

for index in range(0, scale):
ratio = histogram[index] / pixels
brightness += ratio * (-scale + index)

return 1 if brightness == 255 else brightness / scale

def avg_color(path):
img = Image.open(path)
width, height = img.size

r_total = 0
g_total = 0
b_total = 0

count = 0
for x in range(0, width):
for y in range(0, height):
r, g, b = img.getpixel((x,y))
r_total += r
g_total += g
b_total += b
count += 1

return (r_total/count, g_total/count, b_total/count)

def entropy(path):
rgbImg = io.imread(path)
grayImg = img_as_ubyte(color.rgb2gray(rgbImg))
return sklearn_entropy(grayImg)
```

```
In [14]: for i, row in df.iterrows():
name = row['url_thumb'].split('/')[2]
link = row['url_thumb']
path = "data/thumbnails/%s.jpg" % name
urlretrieve(link, path)
clear_output(wait=True)
print("Retrieved %s: %d" % (name, i))
df.at[i, 'blurriness'] = blurriness(path)
print("Computed blurriness %s: %d" % (name, i))
df.at[i, 'brightness'] = brightness(path)
print("Computed brightness %s: %d" % (name, i))
df.at[i, 'avg_color'] = avg_color(path)
print("Computed avg_color %s: %d" % (name, i))
df.at[i, 'entropy'] = entropy(path)
print("Computed entropy %s: %d" % (name, i))
```

```
Retrieved M-QkaJSMR9M: 264
Computed blurriness M-QkaJSMR9M: 264
Computed brightness M-QkaJSMR9M: 264
Computed avg_color M-QkaJSMR9M: 264
Computed entropy M-QkaJSMR9M: 264
```

From "Engagement and Popularity Dynamics of YouTube Videos and Sensitivity to Meta-Data": word count, punctuation count, character count, Google hits (e.g., if the title is entered into the Google search engine how many results are found), and the Sentiment/Subjectivity of the title computed using Vader [30], and TextBlob.

```
In [16]: df.title.sample(10)
```

```
Out[16]: 142    Bullying -- de la culpa a la vergüenza | Paula...
219    Mis alumnos y las calles del pueblo | Miguel Á...
29     Te invito a creer: Manuel Lozano at TEDxRiodel...
254    Transformar la educación en América Latina | G...
31     La puerta equivocada | Adrián Paenza | TEDxRio...
127    Los reflejos emocionales | Hernán Laperuta | T...
50     No todo es puro verso | Cristina Domenech | TE...
131    TEDxBuenosAires - Marcelo Moguilevsky - 04/08/10
168    ¿De qué hablan los delfines? | Marcelo Magnasc...
148    TEDxBuenosAires - Mercedes Salado Puerto - 04/...
Name: title, dtype: object
```

```
In [86]: df['word_count'] = 0
df['punctuation_count'] = 0
df['character_count'] = 0
df['google_hits'] = 0
df['sentiment_vader'] = 0.0
df['sentiment_blob_polarity'] = 0.0
df['sentiment_blob_subjectivity'] = 0.0
```

```
In [92]: def get_hits(searchfor):
    try:
        r = requests.get('http://www.google.com/search',
                           params={'q': '%s' % searchfor,
                                   "tbs": "li:1", # Google doesn't guess for alternat
                                   "hl": 'en'}
                           )
        time.sleep(3) # So google doesn't think we're bots. Wait, but we are...
        soup = BeautifulSoup(r.text, "lxml")
        hits_string = soup.find('div',{'id':'resultStats'}).text
        hits = int(''.join([c for c in hits_string if c.isdigit()]))
    except Exception as e:
        print(e)
        print(searchfor)
        print(hits_string)
        raise e
    if hits == -1:
        print(r)
        print(hits_string)
    return hits
```

```
In [94]: words_series = df.title.str.strip().str.split('[\W_]+')
analyzer = SentimentIntensityAnalyzer()
count = lambda l1,l2: sum([1 for x in l1 if x in l2])
for i, row in df.iterrows():
    df.at[i, 'word_count'] = len(list(filter(None, words_series.loc[i])))
    df.at[i, 'punctuation_count'] = count(row['title'],set(string.punctuation))
    df.at[i, 'character_count'] = len(row['title'])
    if row['google_hits'] <= 0: # Sensitive, let's not recalculate in case google
        df.at[i, 'google_hits'] = get_hits(row['title'])
        df.to_csv("draft_saved_just_in_case.csv")
    df.at[i, 'sentiment_vader'] = analyzer.polarity_scores(row['title'])['compound']
    df.at[i, 'sentiment_blob_polarity'] = TextBlob(row['title']).sentiment.polarit
    df.at[i, 'sentiment_blob_subjectivity'] = TextBlob(row['title']).sentiment.sub
    clear_output(wait=True)
    print(i, row['title'], df.at[i, 'google_hits'])
```

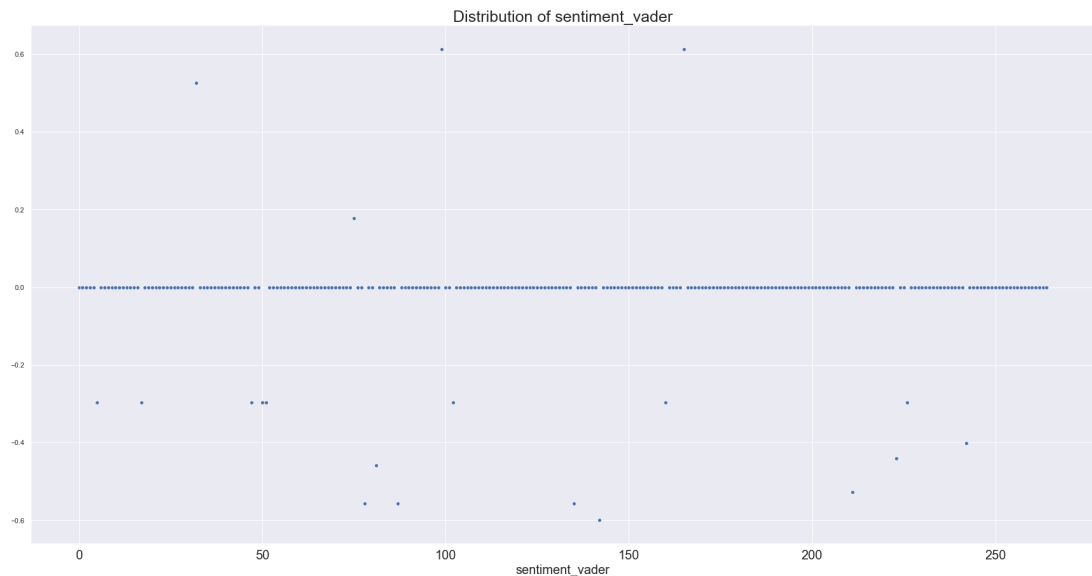
264 Performance | Saxaje | TEDxRiodelaPlataED 8230000000

```
In [98]: # df.to_csv('data/processed_v1.csv')
```

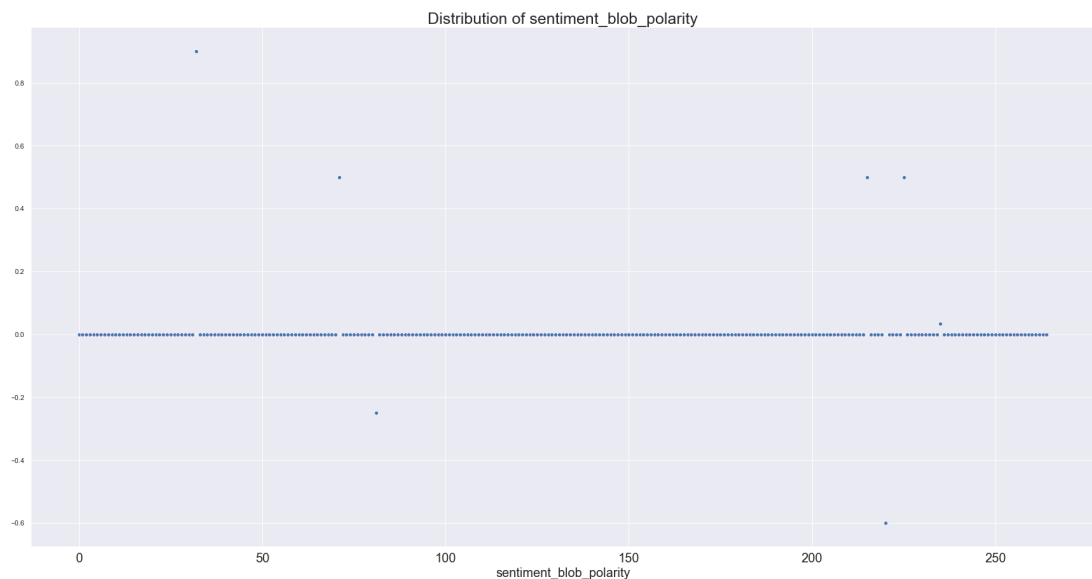
```
In [102]: # df = pd.read_csv("data/processed_v1.csv", index_col=[0])
```

Let's check distribution for the sentiment analysis.

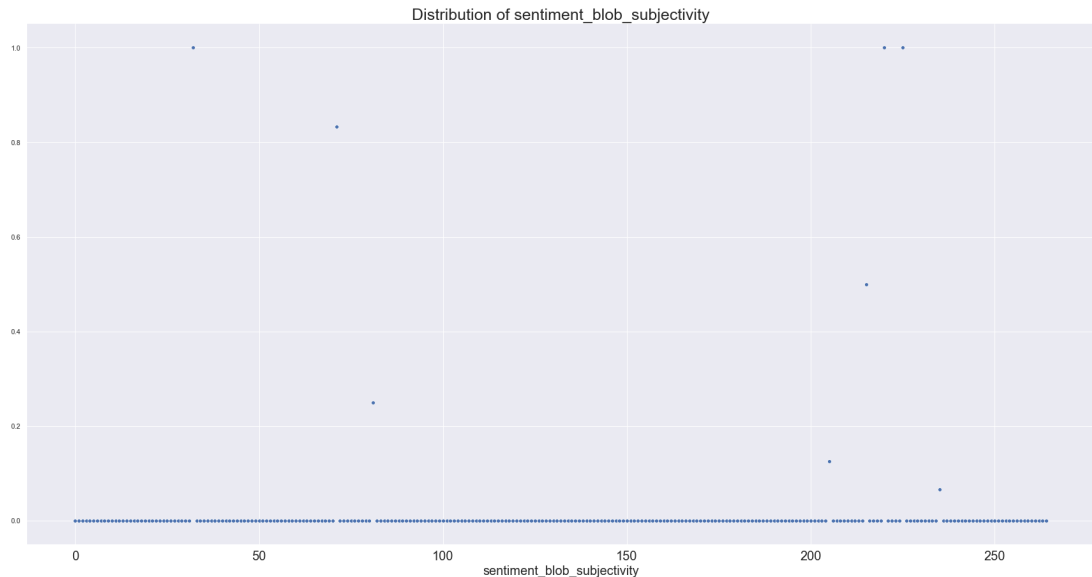
```
In [200]: fig, ax = plt.subplots(figsize=(30,15))
ax = sns.scatterplot(data=df.sentiment_vader)
plt.xlabel('sentiment_vader', fontsize=20)
plt.title('Distribution of sentiment_vader', fontsize=25)
for tick in ax.xaxis.get_major_ticks():
    tick.label.set_fontsize(20)
sns.despine()
plt.plot();
```



```
In [202]: fig, ax = plt.subplots(figsize=(30,15))
ax = sns.scatterplot(data=df.sentiment_blob_polarity)
plt.xlabel('sentiment_blob_polarity', fontsize=20)
plt.title('Distribution of sentiment_blob_polarity', fontsize=25)
for tick in ax.xaxis.get_major_ticks():
    tick.label.set_fontsize(20)
sns.despine()
plt.plot();
```



```
In [203]: fig, ax = plt.subplots(figsize=(30,15))
ax = sns.scatterplot(data=df.sentiment_blob_subjectivity)
plt.xlabel('sentiment_blob_subjectivity', fontsize=20)
plt.title('Distribution of sentiment_blob_subjectivity', fontsize=25)
for tick in ax.xaxis.get_major_ticks():
    tick.label.set_fontsize(20)
sns.despine()
plt.plot();
```



Not very useful... let's leave it at that for the moment

From "Engagement and Popularity Dynamics of YouTube Videos and Sensitivity to Meta-Data":

For the Keywords, seven meta-level features are computed which include: the number of keywords, and keyword length

```
In [185]: df['tags'] = df['tags'].apply(str).apply(lambda x: '' if x=='nan' else x)
```

```
In [191]: df['n_keywords'] = 0
df['keyword_length'] = 0.0
```

```
In [192]: for i, row in df.iterrows():
    if row['tags']:
        tags = row['tags'].split(',')
        df.at[i, 'n_keywords'] = len(tags)
        df.at[i, 'keyword_length'] = sum( map(len, tags) ) / len(tags)
        clear_output(wait=True)
        print(i, row['title'], sum( map(len, tags) ) / len(tags))
```

264 Performance | Saxaje | TEDxRiodelaPlataED 9.4

From "Engagement and Popularity Dynamics of YouTube Videos and Sensitivity to Meta-Data":

In addition, to the above 49 meta-level features, we also include auxiliary user meta-level features including: the number of subscribers, resolution of the thumbnail used, category of the video, the length of the video, and the first day view count of the video.


```
In [195]: df.columns
```

```
Out[195]: Index(['event_name', 'video_id', 'title', 'url_thumb', 'tags', 'view_count',
               'like_count', 'dislike_count', 'comment_count', 'created', 'blurriness',
               'brightness', 'contrast', 'overexposure', 'entropy', 'avg_color',
               'word_count', 'punctuation_count', 'character_count', 'google_hits',
               'sentiment_vader', 'sentiment_blob_polarity',
               'sentiment_blob_subjectivity', 'n_keywords', 'keyword_length'],
              dtype='object')
```

```
In [18]: df['created'] = pd.to_datetime(df['created'])
df['created_day'] = [d.day for d in df['created']]
df['created_month'] = [d.month for d in df['created']]
df['created_year'] = [d.year for d in df['created']]
df['created_hour'] = [d.hour for d in df['created']]
df['created_minute'] = [d.minute for d in df['created']]
df['created_second'] = [d.second for d in df['created']]
```

```
In [20]: #df.to_csv('data/processed_v2.csv')
```

```
In [204]: # df = pd.read_csv("data/processed_v2.csv", index_col=[0])
```

Let's start analyzing and doing stuff. Add features later.

```
In [205]: df = df.drop(['video_id', 'title', 'url_thumb', 'tags', 'created'], axis=1)
```

Split color

```
In [206]: df[['avg_color_r', 'avg_color_g', 'avg_color_b']] = df['avg_color'].apply(eval).ap
```

```
In [207]: df = df.drop(['avg_color'], axis=1)
```

One hot encode event

```
In [208]: df.event_name.unique()
```

```
Out[208]: array(['CanalTEDxRiodelaPlata', 'TEDxRiodelaPlata2013',
               'TEDxJoven@RiodelaPlata 2011', 'TEDxRiodelaPlata2014',
               'TEDxRiodelaPlata2015', 'TEDxRiodelaPlata2011',
               'TEDxRiodelaPlata2017', 'TEDxRiodelaPlataED2015',
               'TEDxRiodelaPlata2016', 'TEDxRiodelaPlata2012', 'TEDxChange',
               'TEDxBuenosAires2010', 'TEDxJoven@RiodelaPlata 2012',
               'TEDxRiodelaPlata2010', 'TEDxRiodelaPlataSalon2016'], dtype=object)
```

```
In [209]: df = pd.get_dummies(df, prefix=['event'], columns=['event_name'])
```

```
In [210]: df.columns = ['view_count', 'like_count', 'dislike_count', 'comment_count',
                        'blurriness', 'brightness', 'contrast', 'overexposure', 'entropy',
                        'word_count', 'punctuation_count', 'character_count', 'google_hits',
                        'sentiment_vader', 'sentiment_blob_polarity',
                        'sentiment_blob_subjectivity', 'n_keywords', 'keyword_length',
                        'created_day', 'created_month', 'created_year', 'created_hour',
                        'created_minute', 'created_second', 'avg_color_r', 'avg_color_g',
                        'avg_color_b', 'event_CanalTEDxRiodelaPlata',
                        'event_TEDxBuenosAires2010', 'event_TEDxChange',
                        'event_TEDxJoven@RiodelaPlata2011',
                        'event_TEDxJoven@RiodelaPlata2012', 'event_TEDxRiodelaPlata2010',
                        'event_TEDxRiodelaPlata2011', 'event_TEDxRiodelaPlata2012',
                        'event_TEDxRiodelaPlata2013', 'event_TEDxRiodelaPlata2014',
                        'event_TEDxRiodelaPlata2015', 'event_TEDxRiodelaPlata2016',
                        'event_TEDxRiodelaPlata2017', 'event_TEDxRiodelaPlataED2015',
                        'event_TEDxRiodelaPlataSalon2016']
```

```
In [211]: y = df.view_count
          X = df[df.columns[1:]]
```

From "Engagement and Popularity Dynamics of YouTube Videos and Sensitivity to Meta-Data":

Models: the machine learning methods we utilize include: the Extreme Learning Machine (ELM), Feed-Forward Neural Network (FFNN), Stacked Auto-Encoder Deep Neural-Network, Elasticnet, Lasso, Relaxed Lasso, Quantile Regression with Lasso, Conditional Inference Random Forest (CIRF), Boosted Generalized Additive Model, Bagged MARS using gCV Pruning, Generalized Linear Model with Stepwise Feature Selection using Akaike information criterion, and Spike and Slab Regression.

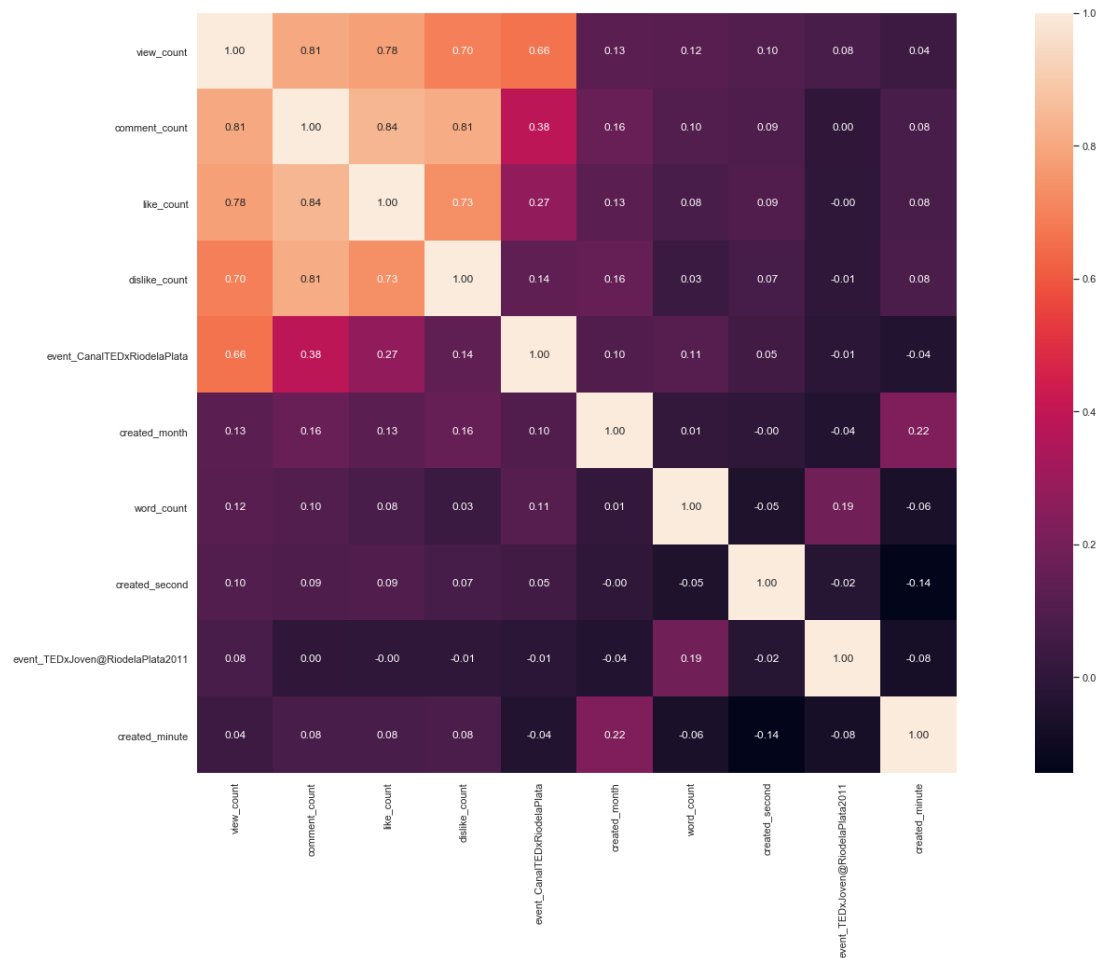
```
In [212]: from sklearn_extensions.extreme_learning_machines.elm import ELMRegressor
          from sklearn.linear_model import ElasticNet, Lasso, LinearRegression
          from sklearn.ensemble import GradientBoostingRegressor
          from pygam import LinearGAM
          from sklearn.ensemble import AdaBoostRegressor
          from xgboost import XGBRegressor
          from sklearn.tree import DecisionTreeRegressor
          from sklearn.neighbors import KNeighborsRegressor

          from sklearn.model_selection import train_test_split, KFold, cross_val_score
```

```
In [213]: correlation = df.corr(method='pearson')
          columns = correlation.nlargest(10, 'view_count').index
          columns
```

```
Out[213]: Index(['view_count', 'comment_count', 'like_count', 'dislike_count',
                 'event_CanalTEDxRiodelaPlata', 'created_month', 'word_count',
                 'created_second', 'event_TEDxJoven@RiodelaPlata2011', 'created_minute'],
                 dtype='object')
```

```
In [214]: correlation_map = np.corrcoef(df[columns].values.T)
fig, ax = plt.subplots(figsize=(30,15))
sns.set(font_scale=1.0)
heatmap = sns.heatmap(correlation_map, cbar=True, annot=True, square=True, fmt='.2')
plt.show()
```



We don't have a lot of correlated features

```
In [215]: df.hist(figsize=(30,30));
```



```
In [216]: X_train, X_test, y_train, y_test = train_test_split (X, y, test_size = 0.20, random_state=42)
```

```
In [218]: from sklearn.preprocessing import StandardScaler
scalerX = StandardScaler()
scalerY = StandardScaler()
X_train = scalerX.fit_transform(X_train)
y_train = scalerY.fit_transform(y_train.reshape(-1, 1))
X_test = scalerX.fit_transform(X_test)
y_test = scalerY.fit_transform(y_test.reshape(-1, 1))
```

Let's check performance of the different models

```

In [220]: models = [
    ("ELMRegressor", ELMRegressor()), ("ElasticNet", ElasticNet()), ("Lasso", Lasso),
    ("LinearRegression", LinearRegression()), ("GradientBoostingRegressor", GradientBoostingRegressor()),
    ("GradientBoostingRegressor w/quantile loss", GradientBoostingRegressor(loss='quantile')),
    ("LinearGAM", LinearGAM()), ("AdaBoostRegressor", AdaBoostRegressor()), ("XGBRegressor", XGBRegressor()),
    ("DecisionTreeRegressor", DecisionTreeRegressor()), ("KNeighborsRegressor", KNeighborsRegressor())
]

import warnings
warnings.filterwarnings('ignore')

results = []
names = []
for name, model in models:
    kfold = KFold(n_splits=10, random_state=21)
    cv_results = cross_val_score(model, X_train, y_train, cv=kfold, scoring='neg_mean_squared_error')
    results.append(cv_results.mean())
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
print("Best model: %s" % str(sorted(list(zip(results, names)), key=lambda tup: tup[0])))

ELMRegressor: -0.889908 (1.012574)
ElasticNet: -0.809974 (1.318992)
Lasso: -1.010317 (1.393977)
LinearRegression: -0.454327 (0.810584)
GradientBoostingRegressor: -0.308878 (0.634281)
GradientBoostingRegressor w/quantile loss: -0.353544 (0.509432)
LinearGAM: -0.760800 (1.079107)
AdaBoostRegressor: -0.420131 (0.634911)
XGBRegressor: -0.312693 (0.680793)
DecisionTreeRegressor: -0.929138 (1.286812)
KNeighborsRegressor: -0.684558 (1.000135)
Best model: (-0.30887841110299236, 'GradientBoostingRegressor')

```

Grid search for hyperparameters:

```
In [222]: from sklearn.model_selection import GridSearchCV

param_grid = dict(n_estimators=np.array([50,100,200,300,400]),
                  loss=['ls', 'lad', 'huber', 'quantile'],
                  learning_rate=np.array([0.1, 0.05, 0.01]),
                  max_depth=np.array([3,5,10]))
model = GradientBoostingRegressor(random_state=21)
kfold = KFold(n_splits=10, random_state=21)
grid = GridSearchCV(estimator=model, param_grid=param_grid, scoring='neg_mean_squa
grid_result = grid.fit(X_train, y_train)

means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
```

```
-0.291833 (0.633911) with: {'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 3,
'n_estimators': 50}
-0.289982 (0.632731) with: {'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 3,
'n_estimators': 100}
-0.289588 (0.632717) with: {'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 3,
'n_estimators': 200}
-0.289528 (0.632892) with: {'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 3,
'n_estimators': 300}
-0.289590 (0.633003) with: {'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 3,
'n_estimators': 400}
-0.603071 (1.126002) with: {'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 5,
'n_estimators': 50}
-0.603566 (1.125133) with: {'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 5,
'n_estimators': 100}
-0.603530 (1.125008) with: {'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 5,
'n_estimators': 200}
-0.603528 (1.125006) with: {'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 5,
'n_estimators': 300}
-0.603528 (1.125006) with: {'learning_rate': 0.1, 'loss': 'ls', 'max_depth': 5,
'n_estimators': 400}
```

So

```
In [223]: model = GradientBoostingRegressor(**grid_result.best_params_)
```

```
In [224]: model.fit(X_train, y_train)
```

```
Out[224]: GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,
learning_rate=0.1, loss='ls', max_depth=3, max_features=None,
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=300, presort='auto', random_state=None,
subsample=1.0, verbose=0, warm_start=False)
```

```
In [225]: from sklearn.metrics import mean_squared_error
y_pred = model.predict(X_test)
print(mean_squared_error(y_test, y_pred))
```

```
0.36232661107619896
```

```
In [226]: compare = pd.DataFrame({'Prediction': y_pred, 'Test Data' : y_test.flatten()})
          compare.head(10)
```

Out[226]:

	Prediction	Test Data
0	-0.349741	-0.529010
1	-0.338426	-0.419116
2	-0.153162	-0.363055
3	2.322078	0.984878
4	1.885398	3.303815
5	-0.370476	-0.475171
6	-0.377834	-0.585905
7	-0.046461	0.097545
8	-0.349484	-0.536913
9	-0.326438	-0.444637

Let's descale

```
In [230]: actual_y_test = scaler.inverse_transform(y_test).flatten()
          actual_predicted = scaler.inverse_transform(y_pred)
          diff = abs(actual_y_test - actual_predicted)

          compare_actual = pd.DataFrame({'Test Data': actual_y_test, 'Predicted views' : actual_predicted})

          compare_actual = compare_actual.astype(int)
          compare_actual[['Test Data', 'Predicted views', 'Difference']].head(10)
```

Out[230]:

	Test Data	Predicted views	Difference
0	9978	37085	27107
1	26595	38796	12201
2	35072	66809	31737
3	238892	441089	202197
4	589537	375059	214477
5	18119	33949	15830
6	1375	32837	31462
7	104719	82943	21775
8	8783	37124	28341
9	22736	40608	17872

```
In [228]: np.set_printoptions(formatter={'float_kind':'{:f}'.format})
important_features_dict = {}
for x,i in zip(X.columns, model.feature_importances_):
    important_features_dict[x]=i

important_features_list = sorted(important_features_dict,
                                key=important_features_dict.get,
                                reverse=True)

for feature in important_features_list:
    print("%s: %f" % (feature, important_features_dict[feature]))

like_count: 0.157566
comment_count: 0.083433
dislike_count: 0.068754
entropy: 0.067336
blurriness: 0.060509
n_keywords: 0.054015
keyword_length: 0.049715
avg_color_r: 0.049344
avg_color_b: 0.048850
avg_color_g: 0.048818
character_count: 0.046719
google_hits: 0.042824
created_second: 0.040307
created_minute: 0.034012
brightness: 0.028513
created_hour: 0.019080
punctuation_count: 0.014174
sentiment_vader: 0.013227
word_count: 0.011739
created_day: 0.008017
created_month: 0.007740
event_TEDxRiodelaPlata2017: 0.006784
event_TEDxRiodelaPlata2012: 0.006389
event_TEDxRiodelaPlata2016: 0.005892
event_TEDxRiodelaPlata2013: 0.004568
event_TEDxBuenosAires2010: 0.003785
event_TEDxRiodelaPlataED2015: 0.003708
sentiment_blob_polarity: 0.003611
event_TEDxJoven@RiodelaPlata2012: 0.002403
event_TEDxJoven@RiodelaPlata2011: 0.002388
event_TEDxRiodelaPlata2011: 0.002077
sentiment_blob_subjectivity: 0.001431
event_TEDxRiodelaPlata2015: 0.000678
event_TEDxChange: 0.000671
event_TEDxRiodelaPlata2014: 0.000612
event_CanalTEDxRiodelaPlata: 0.000313
contrast: 0.000000
overexposure: 0.000000
created_year: 0.000000
event_TEDxRiodelaPlata2010: 0.000000
event_TEDxRiodelaPlataSalon2016: 0.000000
```

Some heavy reengineering in the features need to be done.

In []: